

Lecture 6.1

Transport Layer: Process-to-Process Delivery: UDP and TCP


Dr. Vandana Kushwaha

Department of Computer Science
Institute of Science, BHU, Varanasi

PROCESS-TO-PROCESS DELIVERY

- Communication on the **Internet** is not defined as the exchange of data between two nodes or between two hosts.
- But real communication takes place between two processes (application programs), ie. we need **process-to-process delivery**.
- However, at any moment, several processes may be running on the source host and several on the destination host.
- To complete the **delivery**, we need a mechanism to deliver data from one of these processes running on the source host to the corresponding process running on the destination host.
- The **Transport layer** is responsible for **process-to-process delivery**-the delivery of a packet, part of a message, from one process to another.

PROCESS-TO-PROCESS DELIVERY



Client/Server Paradigm


- Although there are **several ways** to achieve **process-to-process** communication, the **most common** one is through the **client/server paradigm**.
- A **process** on the **local host**, called a **client**, needs **services** from a **process** usually on the **remote host**, called a **server**.
- **Both processes** (**client** and **server**) have the **same name**.
- For **example**, to open a **web page** from a **remote machine(Server)**, we need a **HTTP client** process running on the **local host(Client)** and a **HTTP server** process running on a **remote machine(Server)**.
- A **remote computer** can run **several server programs** at the **same time**, just as **local computers** can run one or more **client programs** at the **same time**.

Port Address

- In the **Internet model**, the **port numbers** are **16-bit integers** between **0** and **65,535**.
- The **client program** defines **itself** with a **port number**, chosen **randomly** by the **transport layer software** running on the **client host**.
- This is the **ephemeral port number**.
- The **server process** must also **define itself** with a **port number**.
- This **port number**, however, **cannot be chosen randomly**.
- The **Internet** has **decided** to use **universal port numbers** for **servers**; these are called **well-known port numbers**.


Types of Port Numbers

- The **IANA** (Internet Assigned Number Authority) has divided the **port numbers** into **three ranges**: *well known*, *registered*, and *dynamic (or private)*.
- i. **Well-known ports.** The ports ranging from **0 to 1023** are assigned and controlled by **IANA**. These are the well-known ports.
- ii. **Registered ports.** The ports ranging from **1024 to 49,151** are not assigned or controlled by **IANA**. They can only be registered with **IANA** to prevent duplication. These are used by **vendors** for their **own server applications**.
- iii. **Dynamic ports.** The ports ranging from **49,152 to 65,535** are neither controlled nor registered. They can be used by any **process**. These are the **ephemeral ports**.



Socket Addresses

- **Process-to-process delivery** needs **two identifiers**, IP address and the port number, at each end to make a **connection**.
- The **combination** of an **IP address** and a **port number** is called a **socket address**.
- The **client socket address** defines the **client process uniquely** just as the **server socket address** defines the **server process uniquely**.



Connectionless Versus Connection-Oriented Service

- A **transport layer protocol** can either be **connectionless** or **connection-oriented**.
- ***Connectionless Service***
- In a **connectionless service**, the **packets** are sent from one party to another with no need for connection establishment or connection release.
- The **packets** are **not numbered**; they may be **delayed** or **lost** or may arrive **out of sequence**.
- There is **no acknowledgment** either.
- ***Connection Oriented Service***
- In a **connection-oriented service**, a **connection** is **first established** between the **sender** and the **receiver**, **data are transferred**.
- At the **end**, the **connection** is **released**.

Reliable Versus Unreliable

- The **Transport layer service** can be **reliable** or **unreliable**.
- If the **application layer program** needs **reliability**, we use a **reliable transport layer protocol** by implementing **flow** and **error control** at the **transport layer**.
- This means a **slower** and more **complex service**.
- **TCP** is a **connection oriented** and **reliable delivery protocol**.
- On the other hand, if the **application program** does not need **reliability** because
 - It uses its **own flow** and **error control** mechanism or
 - It **needs fast service** or
 - The **nature** of the **service** does not demand **flow** and **error control** (real-time applications),
- Then an **unreliable protocol** can be used.
- **UDP** is **connectionless** and **unreliable delivery protocol**.

Protocols at Transport Layer

The original **TCP/IP protocol suite** specifies **two protocols** for the **transport layer**: **UDP** and **TCP**.

USER DATAGRAM PROTOCOL (UDP)

- The *User Datagram Protocol (UDP)* is called a **connectionless, unreliable transport protocol**.
- It **does not add** anything to the **services** of **IP** except to provide **process-to process communication** instead of **host-to-host** communication.
- Also, it performs **very limited** error checking.
- **UDP** is a **very simple protocol** using a **minimum of overhead**.
- If a **process** wants to send a **small message** and does **not care much** about **reliability**, it can use **UDP**.
- Sending a **small message** by using **UDP** takes much **less interaction** between the **sender** and **receiver** than using **TCP**.


Format of User Datagram

UDP packets, called user datagrams, have a fixed-size header of 8 bytes. The fields are as follows:

i. Source port number.

- This is the port number used by the process running on the source host.
- It is 16 bits long, which means that the port number can range from 0 to 65,535.
- If the source host is the client (a client sending a request), the port number, in most cases, is an ephemeral port number requested by the process and chosen by the UDP software running on the source host.
- If the source host is the server (a server sending a response), the port number, in most cases, is a well-known port number.

Format of User Datagram



ii. Destination port number.

- This is the **16 bits** long **port number** used by the **process** running on the **destination host**.
- If the **destination host** is the **server** (a client sending a request), the port number, in most cases, is a **well-known port number**.
- If the **destination host** is the **client** (a server sending a response), the port number, in most cases, is an **ephemeral port number**. In this case, the **server** copies the **ephemeral port number** it has received in **the request packet**.

Format of User Datagram

iii. Length.

- This is a **16-bit field** that defines the **total length** of the **user datagram, header plus data**.
- The **16 bits** can define a total length of **0 to 65,535 bytes**.
- However, the **total length** needs to be **less** because a **UDP user datagram** is stored in an **IP datagram** with a **total length of 65,535 bytes**.
- The **length field** in a **UDP user datagram** is actually **not necessary**.
- A **user datagram** is **encapsulated** in an **IP datagram**.
 - **UDP length = IP length - IP header's length**

iv. Checksum.

This field is used to **detect errors** over the entire user datagram (header plus data).

UDP Operation

Connectionless Services

- As mentioned previously, **UDP** provides a **connectionless service**.
- This means that **each user datagram** sent by **UDP** is an **independent datagram**.
- There is **no relationship** between the different **user datagrams** even if they are coming from the **same source process** and going to the **same destination program**.
- The **user datagrams** are **not numbered**.
- Also, there is **no connection establishment** and **no connection termination**, as is the case for **TCP**.
- This means that **each user datagram** can travel on a **different path**.

UDP Operation


Flow and Error Control

- UDP is a **very simple, unreliable** transport protocol.
- There is **no flow control** and hence no **window mechanism**.
- The **receiver may overflow** with **incoming messages**.
- There is **no error control** mechanism in **UDP** except for the **checksum**.
- This means that the **sender does not know** if a message has been **lost** or **duplicated**.
- When the **receiver detects an error** through the **checksum**, the **user datagram is silently discarded**.

Encapsulation and Decapsulation

- To send a **message** from one **process** to another, the **UDP** protocol **encapsulates** and **decapsulates** messages in an **IP datagram**.

UDP Encapsulation and Decapsulation Process



Use of UDP

- **UDP** is suitable for a process that requires simple **request-response communication** with **little concern** for **flow** and **error control**.
- It is **not** usually **used** for a **process** such as **FTP (File Transfer Protocol)** that needs to send **bulk data**.
- **UDP** is suitable for a **process** with **internal flow** and **error control mechanisms**. For example, the **Trivial File Transfer Protocol (TFTP)** process includes **flow** and **error control**.
- **UDP** is a suitable **transport protocol** for **multicasting**. Multicasting capability is **embedded** in the **UDP** software but **not in** the **TCP** software.
- **UDP** is used for **management processes** such as **SNMP (Simple Network Management Protocol)**.
- **UDP** is used for some **route updating protocols** such as **Routing Information Protocol (RIP)**.

TCP (Transmission Control Protocol)


- **TCP (Transmission Control Protocol)** is a *connection-oriented, reliable* transport protocol.
- It adds **connection-oriented** and **reliability** features to the **services of IP**.
- In addition, **TCP** uses **flow** and **error control** mechanisms at the **transport level**.
- **Services offered by TCP**
 - i. **Process-to-Process Communication**
 - Like **UDP**, **TCP** provides **process-to-process communication** using **port numbers**.

TCP Services

ii. Stream Delivery Service

- TCP, allows the **sending process** to **deliver data** as a **stream of bytes** and allows the **receiving process** to **obtain data** as a **stream of bytes**.
- TCP creates an **environment** in which the **two processes** seem to be **connected** by an **imaginary "tube"** that carries their **data** across the **Internet**.
- The **sending process** **produces** (writes to) the **stream of bytes**, and the **receiving process** **consumes** (reads from) them.
- The **sending** and the **receiving processes** may not **write** or **read data** at the **same speed**.
- There are **two buffers**, the **sending buffer** and the **receiving buffer**, one for each direction.

TCP Services: Stream Delivery Service



TCP Services

iii. Segments

- TCP **groups** a number of bytes together into a **packet** called a **segment**.
- TCP adds a **header** to each **segment** (for control purposes) and **delivers** the **segment** to the **IP layer** for transmission.
- The **segments** are **encapsulated** in IP datagrams and transmitted. The **segments** are **not necessarily the same size**.

iv. Full-Duplex Communication

- TCP offers **full-duplex service**, in which **data** can **flow in both directions** at the **same time**.
- Each **TCP** then has a **sending and receiving buffer**, and **segments** move in **both directions**.

Connection-Oriented & Reliable Service

- TCP, unlike UDP, is a **connection-oriented protocol**.
- When a **process** at **site A** wants to **send and receive data** from another process at **site B**, the following **three events** occur:
 1. The **two TCPs establish a connection** between them.
 2. **Data are exchanged** in both directions.
 3. The **connection is terminated**.
- TCP is a **reliable** transport protocol.
- It uses an **acknowledgment mechanism** to check the **safe and sound arrival** of **data**.

TCP Features

A. Numbering System

TCP software keeps track of the segments being transmitted or received using two fields called the sequence number and the acknowledgment number.

i. Byte Number

- TCP numbers all data bytes that are transmitted in a connection.
- Numbering is independent in each direction.
- When TCP receives bytes of data from a process, it stores them in the sending buffer and numbers them.
- The numbering does not necessarily start from 0.
- Instead, TCP generates a random number between 0 and $2^{32} - 1$ for the number of the first byte.
- For example, if the random number happens to be 1057 and the total data to be sent are 6000 bytes, the bytes are numbered from 1057 to 7056.

TCP Features

ii. Sequence Number

- After the **bytes** have been **numbered**, **TCP** assigns a **sequence number** to each segment that is being sent.
- *The **sequence number** for each segment is the **number of the first byte carried in that segment**.*

Example : Suppose a TCP connection is transferring a file of 5000 bytes. The first byte is numbered 10,001. What are the sequence numbers for each segment if data are sent in five segments, each carrying 1000 bytes?

Solution

The following shows the sequence number for each segment:

- Segment 1 Sequence Number: 10,001 (range: 10,001 to 11,000)
- Segment 2 Sequence Number: 11,001 (range: 11,001 to 12,000)
- Segment 3 Sequence Number: 12,001 (range: 12,001 to 13,000)
- Segment 4 Sequence Number: 13,001 (range: 13,001 to 14,000)
- Segment 5 Sequence Number: 14,001 (range: 14,001 to 15,000)

TCP Features

iii. Acknowledgment Number

- The **acknowledgment number** defines the **number** of the **next byte** that the receiver **expects** to receive.
- In addition, the **acknowledgment number** is **cumulative**, which means that the receiver takes the number of the **last byte** that it has **received**, safe and sound, **adds 1 to it**, and **announces** this **sum** as the **acknowledgment number**.
- The term **cumulative** here means that if a receiver uses **5643** as an **acknowledgment number**, it has **received** all **bytes** from the **beginning** up to **5642**.
- Note that this does not mean that the **receiver** has **received 5642 bytes** because the **first byte number** does not have to start from **0**.

TCP Features

B. Flow Control

- TCP, unlike UDP, provides ***flow control***.
- The **receiver** of the data controls the **amount of data** that are to be sent by the sender.
- This is done to prevent the receiver from being **overwhelmed with data**.
- The numbering system allows TCP to use a **byte-oriented flow control**.

C. Error Control


- To provide **reliable service**, **TCP** implements an **error control mechanism**.
- Although **error control** considers a **segment** as the **unit of data** for error detection (loss or corrupted segments), error control is **byte-oriented**.

D. Congestion Control

- **TCP**, unlike **UDP**, takes into account **congestion** in the **network**.
- The **amount of data** sent by a **sender** is not only controlled by the **receiver** (flow control), but is also **determined** by the **level of congestion** in the **network**.

TCP Segment Format

- The **TCP segment** consists of a **20 to 60 byte header**, followed by **data** from the application program.
- The **header** is **20 bytes** if there are **no options** and up to **60 bytes** if it contains **options**.



TCP Segment Format

Source port address.

- This is a **16-bit field** that defines the **port number** of the **application program** in the host that is **sending** the segment.

Destination port address.

- This is a **16-bit field** that defines the **port number** of the **application program** in the host that is **receiving** the segment.

Sequence number.

- This **32-bit field** defines the **number** assigned to the **first byte** of **data** contained in this **segment**.
- During **connection establishment**, each party uses a **random number generator** to create an **initial sequence number** (ISN), which is usually different in each direction.

TCP Segment Format

Acknowledgment number.

- This **32-bit field** defines the **byte number** that the **receiver** of the **segment** is **expecting to receive** from the other party.
- If the **receiver** of the **segment** has **successfully received** **byte number** x from the other party, it defines $x + 1$ as the **acknowledgment number**.
- **Acknowledgment** and **data** can be **piggybacked** together.

Header length.

- This **4-bit field** indicates the number of **4-byte words** in the **TCP header**.
- The **length** of the **header** can be between **20 and 60 bytes**.
- Therefore, the **value** of this **field** can be **between** 5 ($5 \times 4 = 20$) and 15 ($15 \times 4 = 60$).

Reserved.

- This is a **6-bit field reserved for future use**.

TCP Segment Format

Control Flags

- This field defines **6 different control bits** or **flags** as shown in below.
- **One or more** of these **bits** can be **set** at a time.

<i>Flag</i>	<i>Description</i>
URG	The value of the urgent pointer field is valid.
ACK	The value of the acknowledgment field is valid.
PSH	Push the data.
RST	Reset the connection.
SYN	Synchronize sequence numbers during connection.
FIN	Terminate the connection.

TCP Segment Format

Window size

- This **field** defines the **size of the window**, in **bytes**, that the **other party must maintain**.
- Note that the **length** of this **field** is **16 bits**, which means that the **maximum size** of the **window** is **65,535 bytes**.
- This value is normally referred to as the **receiving window (rwnd)** and is **determined by the receiver**.
- The **sender must obey** the **dictation** of the **receiver** in this case.

TCP Segment Format

Checksum

- This **16-bit field** contains the checksum for **header+data security**.

Urgent pointer

- This **16-bit field**, which is **valid only** if the **urgent flag** is **set**, is used when the **segment** contains **urgent data**.
- It **defines** the **number** that must be **added** to the **sequence number** to obtain the **number of the last urgent byte** in the **data section** of the **segment**.

Options

- There can be up to **40 bytes** of **optional information** in the **TCP header**.

A TCP Connection

- TCP is connection-oriented protocol.
- A connection-oriented transport protocol establishes a virtual path between the source and destination.
- All the segments belonging to a message are then sent over this virtual path.
- Using a single virtual pathway for the entire message facilitates the acknowledgment process as well as retransmission of damaged or lost frames.
- In TCP, connection-oriented transmission requires three phases:
 1. *Connection establishment,*
 2. *Data transfer,*
 3. *Connection termination.*


1. Connection Establishment

- TCP transmits data in **full-duplex mode**.
- When two **TCPs** in **two machines** are **connected**, they are **able to send segments** to each other **simultaneously**.
- This implies that **each party** must **initialize** communication and **get approval** from the **other party before** any **data** are transferred.
- The **connection establishment** in **TCP** is called **three-way handshaking**.
- For **example**, an **application program**, called the **client**, wants to make a **connection** with another **application program**, called the **server**, using **TCP** as the **transport layer protocol**.
- The **process starts** with the **server**.
- The **server program** tells its **TCP** that it is **ready** to **accept** a connection.

Connection Establishment

- This is called a **request** for a ***passive open***.
- Although the **server TCP** is **ready** to **accept** any **connection** from **any machine** in the world, it cannot make the connection itself.
- The **client program** issues a **request** for an ***active open***.
- A **client** that **wishes** to **connect** to an open **server** tells **its TCP** that it **needs** to be **connected** to that **particular server**.
- **TCP** can now **start** the **three-way handshaking process** as shown in **Figure** on next slide.
- To show the process, we use **two time lines**: one at each site.
- Each **segment** has **value** for all its **header fields**.

Connection Establishment using three-way handshaking



Connection establishment using three-way handshaking

- The **three steps** in this **phase** are as follows:

Step 1:

- The **client** sends the **first segment**, a **SYN segment**, in which only the **SYN flag** is **set**.
- This **segment** is for **synchronization** of **sequence numbers**.
- A **SYN segment consumes one sequence number**.
- When the **data transfer** starts, the **sequence number** is **incremented by 1**.
- Thus a **SYN segment cannot carry data**, but it **consumes** one **sequence number**.

Connection establishment using three-way handshaking

Step 2:

- The **Server** sends the **second segment**, a **SYN +ACK** segment, with **2 flag bits set: SYN and ACK**.
- This **segment** has a **dual purpose**.
- It is a **SYN segment** for **communication** in the other direction and serves as the **acknowledgment** for the **SYN segment**.
- It **consumes** one sequence number.
- A **SYN +ACK** segment cannot carry data, but does consume one sequence number.

Connection establishment using three-way handshaking

Step 3:

- The **client** sends the **third segment**.
- This is **just an ACK segment**.
- It **acknowledges** the receipt of the **second segment** with the **ACK flag** and **acknowledgment number field**.
- Note that the **sequence number** in this **segment** is the **same** as the one in the **SYN segment**;
- *The ACK segment does not consume any sequence numbers.*
- *An ACK segment, if carrying no data, consumes no sequence number.*

Data Transfer


- After **connection is established**, bidirectional data transfer can take place.
- The **client** and **server** can **both send data and acknowledgments**.
- The **acknowledgment** is **piggybacked** with the **data**.
- **Figure** on next slide shows an **example**. In this **example**, after **connection is established** (not shown in the figure), the **client** sends **2000 bytes** of data in **two segments**.
 - The **server** then sends **2000 bytes** in **one segment**.
 - The **client** sends **one more** segment.
 - The **first three segments** carry **both data and acknowledgment**, but the **last segment** carries **only an acknowledgment** because there are **no more data** to be sent.

Data Transfer

- Note the **values** of the **sequence** and **acknowledgment numbers**.
- The **data segments** sent by the **client** have the **PSH (push) flag set** so that the **server TCP** knows to **deliver data** to the **server process** as soon as they are **received**.
- The **segment** from the **server**, on the **other hand**, does **not set** the **push flag**.

Data Transfer

The Client sends two 2000 bytes of data in 2 segments at the same time acknowledged previous sending by server: 15001



Pushing Data

- The **sending TCP** uses a **buffer** to store the **stream of data** coming from the **sending application program**.
- The **receiving TCP** also **buffers** the data when they arrive and delivers them to the **application program** when the **application program** is **ready**.
- This type of **flexibility** increases the **efficiency** of TCP.
- However, on occasion the **application program** has **no need** for this **flexibility**.
- For **example**, consider an **application program** that communicates **interactively** with another **application program** on the other end.
- The **application program** on **one site** wants to **send** a **keystroke** to the **application** at the **other site** and want to **receive** an **immediate response**.
- **Delayed transmission** and **delayed delivery** of data may not be acceptable by the **application program**.

Pushing Data

- TCP can handle such a situation.
- The application program at the sending site can request a *push* operation.
- This means that the sending TCP must not wait for the window to be filled.
- It must create a segment and send it immediately.
- The sending TCP must also set the push bit (PSH) to let the receiving TCP know that the segment includes data that must be delivered to the receiving application program as soon as possible and not to wait for more data to come.

Urgent Data

- At some occasion an application program needs to send *urgent bytes*.
- This means that the sending application program wants a **piece of data** to be **read out of order** by the receiving application program.
- As an **example**, suppose that the sending application program is **sending data** to be **processed** by the receiving application program.
- When the **result** of processing **comes back**, the sending application program finds that **everything is wrong**.
- It wants to **abort** the process, but it has **already sent** a **huge amount** of data.
- If it issues an **abort command** (control +C), these two characters will be **stored** at the **end of the receiving TCP buffer**.
- It will be **delivered** to the receiving application program after **all the data have been processed**.

Urgent Data

- The **solution** is to send a segment with the **URG bit set**.
- The **sending application program** tells the **sending TCP** that the **piece of data** is **urgent**.
- The **sending TCP** creates a **segment** and **inserts the urgent data at the beginning of the segment**.
- The **rest of the segment** can contain **normal data** from the buffer.
- The **urgent pointer field in the header** defines the **end of the urgent data** and the **start of normal data**.
- When the **receiving TCP** receives a **segment** with the **URG bit set**, it **extracts the urgent data** from the **segment**, using the **value** of the **urgent pointer**, and **delivers them, out of order**, to the **receiving application program**.

URG vs PSH

- **PSH**
 - Send this **message** to the **application right now — don't wait**.
 - **Flushes** the **TCP buffer** and **pushes data** to the **application**.
- **URG**
 - This **part** of the **message** is **important** — **handle it first**.
 - May **bypass** normal data flow for **specific bytes**.

Connection Termination in TCP

- Any of the **two parties** involved in **exchanging data** (client or server) can **close** the **connection**, although it is **usually initiated by the client**.
- Most implementations today allow **two options** for **connection termination**:
 - i. ***Three-way handshaking*** for **full-close**
 - ii. ***Four-way handshaking*** for **half-close**
- Most **implementations** today allow ***three-way handshaking*** for connection termination.

Connection Termination in TCP

Step 1

- In a **normal situation**, the **client TCP**, after **receiving a close command from the client process**, **sends the first segment**, a **FIN segment** in which the **FIN flag is set**.
- Note that a **FIN segment** can **include the last chunk of data** sent by the **client**, or it can be just a **control segment**.
- If it is **only a control segment**, it **consumes only one sequence number**.

Connection Termination

Step 2

- The **server TCP**, after **receiving** the **FIN segment**, informs its **process** of the **situation** and **sends** the **second segment**, a **FIN +ACK segment**, to **confirm** the **receipt** of the **FIN segment** from the **client** and at the same time to **announce** the **closing of the connection** in the **other direction**.
- This **segment** can also **contain** the **last chunk** of data from the **server**.
- If it **does not carry data**, it **consumes** only **one sequence number**.


Connection Termination

Step 3

- The **client TCP** sends the **last segment**, an **ACK** segment, to **confirm** the receipt of the **FIN segment** from the **TCP server**.
- This **segment** contains the **acknowledgment number**, which is **1 plus** the **sequence number received** in the **FIN segment** from the **server**.
- This **segment cannot carry data** and **consumes no sequence numbers**.

Connection termination using three-way handshaking

The Client TCP, after receive a close command from the client process, sends the 1st segment, a FIN, to initiate active close.



Half-Close

- In **TCP**, *one end can stop sending data while still receiving data*. This is called a **half-close**.
- Although **either end** can issue a **half-close**, it is normally **initiated by the client**.
- It can **occur** when the **server** needs **all the data before processing** can begin. A good **example** is **sorting**.
- When the **client** sends **data** to the **server** to be **sorted**, the **server** needs to receive **all the data** before **sorting** can start.
- This means the **client**, after **sending** all the **data**, can **close the connection** in the **outbound direction**.
- However, **the inbound direction must remain open** to receive the **sorted data**.
- The **server**, after **receiving** the **data**, still **needs time** for **sorting**; its **outbound direction** must remain **open**.


Half-Close

- The **client half-closes** the **connection** by sending a **FIN segment**.
- The **server** accepts the **half-close** by sending the **ACK segment**.
- The **data transfer** from the **client** to the **server stops**.
- The **server**, however, **can still send data**.
- When the **server** has **sent all** the **processed data**, it sends a **FIN segment**, which is **acknowledged** by an **ACK** from the **client**.
- After **half-closing** of the connection, **data can travel from the server to the client** and **acknowledgments can travel from the client to the server**.

Half-Close

- The **client** cannot send any more data to the **server**.
- Although the **client** has received **sequence number $y - 1$** and is **expecting y** , the **server sequence number** is still **$y - 1$** .
- When the **connection** finally **closes**, the **sequence number** of the **last ACK segment** is **still x** , because **no sequence numbers** are **consumed** during **data transfer** in that **direction**.


Half-Close



Flow Control


- TCP uses a **sliding window mechanism**, to handle **flow control**.
- The **Sliding window protocol** used by TCP, however, is something **between** the **Go-Back-N** and **Selective Repeat** **sliding window mechanisms**.
- The **sliding window protocol** in TCP looks like the **Go-Back-N protocol** because it **does not use NAKs**; it looks like **Selective Repeat** because the **receiver holds the out-of-order segments** until the **missing ones arrive**.
- There are **two big differences** between this **sliding window** and the one we used at the **data link layer**.
- First, the **sliding window of TCP is byte-oriented**; the **sliding window** in the **data link layer is frame-oriented**.
- Second, the TCP's **sliding window is of variable size**; the **sliding window** in the **data link layer** was of **fixed size**.

Go-Back-N at Data Link Layer




Go-Back-N Protocol

Selective-Repeat at Data Link Layer



TCP: Flow Control

- The **window spans a portion of the buffer** containing **bytes received** from the **process**.
- The **bytes inside the window** are the **bytes that can be in transit**; they **can be sent** without worrying about **acknowledgment**.
- The **imaginary window** has **two walls**: one **left** and one **right**.
- The **window** is **opened, closed, or shrunk**. These **three activities**, are in the **control of the receiver** (and **depend on congestion** in the **network**), **not the sender**.



Flow Control

- **Opening a window**
 - Means **moving the right wall to the right**.
 - This allows **more new bytes** in the **buffer** that are **eligible** for **sending**.
- **Closing the window**
 - Means **moving the left wall to the right**.
 - This means that **some bytes** have been **acknowledged** and the **sender** need not worry about them anymore.
- **Shrinking the window**
 - Means **moving the right wall to the left**.
 - It means **revoking** the **eligibility** of **some bytes** for **sending**.


Flow Control

- The **size of the window** at one end is determined by the **lesser of two** values:
 - *receiver window (rwnd)*
 - *congestion window (cwnd)*.
- The **receiver window** is the value **advertised** by the **opposite end** in a **segment** containing **acknowledgment**.
- It is the **number of bytes** the **other end** can **accept** before its **buffer overflows** and **data** are **discarded**.
- The **congestion window** is a value **determined** by the **network** to **avoid congestion**.

Flow Control

Example of a Sliding window.

- The **sender** has sent **bytes** up to **202**. Let **cwnd** is **20**.
- The **receiver** has sent an **acknowledgment number** of **200** with an **rwnd** of **9 bytes**.
- The **size of the sender window** is the **minimum of rwnd and cwnd**, or **9 bytes**.
- Bytes **200** to **202** are **sent**, but **not acknowledged**.
- Bytes **203** to **208** **can be sent** without worrying about **acknowledgment**.
- Bytes **209** and **above** **cannot be sent**.



Error Control in TCP

- TCP is a **reliable** transport layer protocol.
- This **means** that an **application program** that **delivers** a **stream of data** to **TCP** relies on TCP to **deliver** the **entire stream** to the **application program** on the other end **in order, without error**, and **without any part lost or duplicated**.
- TCP provides **reliability** using **error control**.
- **Error control** includes **mechanisms** for **detecting corrupted segments, lost segments, out-of-order segments**, and **duplicated segments**.
- **Error control** also includes a **mechanism for correcting errors** after they are detected.
- **Error detection** and **correction** in TCP is achieved through the use of **three simple tools**:
 - **Checksum, Acknowledgment** and **Retransmission**.

Error Control in TCP

1. Checksum

- Each **segment** includes a **checksum field** which is used to check for a **corrupted segment**.
- If the **segment** is **corrupted**, it is **discarded** by the **destination TCP** and is considered as **lost**.
- TCP uses a **16-bit checksum** that is mandatory in every segment.

2. Acknowledgment

- TCP uses **acknowledgments** to **confirm the receipt of data segments**.
- **Control segments** that carry **no data** but **consume a sequence number** are also **acknowledged**.
- **ACK segments** are **never acknowledged**.
- **ACK segments** do not **consume sequence numbers** and are **not acknowledged**.

Error Control in TCP

3. Retransmission

- The **heart** of the **error control mechanism** is the **retransmission of segments**.
- When a **segment** is **corrupted, lost, or delayed**, it is **retransmitted**.
- In **modern implementations**, a **segment** is **retransmitted on two occasions**:
 - when a **Retransmission timer expires** or
 - when the **Sender receives three duplicate ACKs**.
- Note that **no retransmission** occurs for **segments** that do **not consume sequence numbers**.
- In particular, there is **no retransmission** for an **ACK segment**.

Retransmission After RTO

- A recent implementation of TCP maintains one retransmission time-out (RTO) timer for all outstanding (sent, but not acknowledged) segments.
- When the timer matures, the earliest outstanding segment is retransmitted .
- Note that no time-out timer is set for a segment that carries only an acknowledgment, which means that no such segment is resent.
- The value of RTO is dynamic in TCP and is updated based on the round-trip time (RTT) of segments.
- An RTT is the time needed for a segment to reach a destination and for an acknowledgment to be received.

Retransmission After Three Duplicate ACK Segments

- The previous **rule** about **retransmission** of a **segment** is **sufficient** if the **value of RTO is not very large**.
- Sometimes, however, **one segment is lost** and the **receiver receives so many out-of-order segments** that they **cannot be saved** (**limited buffer size**).
- To **alleviate** this **situation**, most implementations today follow the **three-duplicate-ACKs rule** and **retransmit the missing segment immediately**.
- This **feature** is referred to as **Fast Retransmission**.

Out-of-Order Segments


- When a **segment** is **delayed, lost, or discarded**, the **segments** following that segment **arrive out of order**.
- Originally, **TCP** was designed to **discard all out-of-order segments**, resulting in the **retransmission** of the **missing segment** and the **following segments**.
- Most implementations today **do not discard** the **out-of-order segments**.
- They **store** them **temporarily** and **flag** them as **out-of-order segments** until the **missing segment arrives**.
- Note, however, that the **out-of-order segments** are **not delivered** to the **process**.
- **TCP guarantees** that **data** are **delivered** to the **process** **in order**.

Some Scenarios of TCP Operation:

Normal Operation

- The client TCP sends **one segment**; the server TCP sends **three**.
- There are **data to be sent**, so the **segment** displays the **next byte** expected.
- When the **client receives the first segment** from the **server**, it does not have any more data to send; it sends only an **ACK segment**.
- However, the **acknowledgment** needs to be **delayed** for **500 ms** to see if any more segments arrive.
- When the **timer matures**, it **triggers an acknowledgment**.
- When the **next segment** arrives, another **acknowledgment** timer is set.
- However, before it matures, the **third segment** arrives.
- The **arrival** of the **third segment** triggers another acknowledgment(**cumulative ACK**).

TCP: Normal Operation




Lost Segment

- A **lost segment** and a **corrupted segment** are **treated the same way** by the **receiver**.
- A **lost segment** is **discarded** somewhere in the network; a **corrupted segment** is **discarded** by the **receiver itself**. Both are considered **lost**.
- We are assuming that **data transfer** is **unidirectional**: one site is sending, the other is receiving.
- In our scenario, the **sender** sends **segments 1** and **2**, which are **acknowledged immediately** by an **ACK**.
- **Segment 3**, however, is **lost**.
- The **receiver** receives **segment 4**, which is **out of order**.
- The **receiver stores** the **data** in the **segment** in its **buffer** but **leaves a gap** to indicate that there is **no continuity** in the **data**.

Lost Segment

- The **receiver** immediately **sends** an **acknowledgment** to the **sender**, displaying the **next byte it expects**.
- Note that the **receiver stores bytes 801 to 900**, but **never delivers these bytes** to the **application** until the **gap is filled**.
- We have shown the **timer** for the **earliest outstanding segment**.
- The **timer** for this definitely **runs out** because the **receiver never sends** an **acknowledgment** for **lost or out-of-order segments**.
- When the **timer matures**, the **sending TCP resends segment 3**, which **arrives this time** and is **acknowledged** properly.
- Note that the **value** in the **second and third acknowledgments** differs according to the corresponding rule.

TCP: Lost Segment



Fast Retransmission

- This **scenario** is used when the **RTO has a higher value**.
- When the **receiver** receives the **fourth, fifth, and sixth segments**, it **triggers** an **acknowledgment**.
- The **sender** receives **four acknowledgments** with the **same value** (three duplicates).
- Although the **timer for segment 3 has not matured** yet, the **fast transmission** requires that **segment 3**, the segment that is expected by all these acknowledgments, be **resent immediately**.
- Note that only **one segment** is **retransmitted** although **four segments** are **not acknowledged**.
- When the **sender** receives the **retransmitted ACK**, it knows that the **four segments** are **safe and sound** because **acknowledgment is cumulative**.

Fast Retransmission

