

Lecture 6.2

Transport Layer: Congestion Control

Dr. Vandana Kushwaha

Department of Computer Science
Institute of Science, BHU, Varanasi

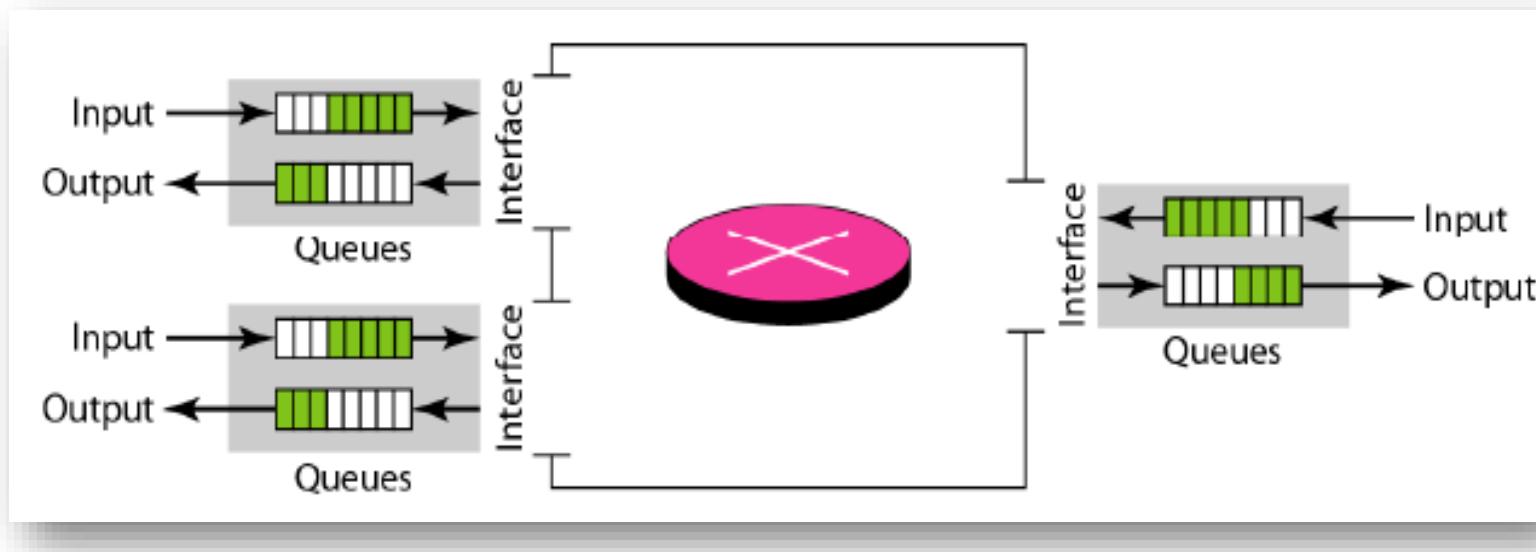
Network Congestion

- Network congestion is one of the serious issues in communication network.
- There are several side effects of network congestion which severely affects the network performance and quality of service.
- Congestion is an important issue in a packet-switched network.
- Congestion in a network may occur if the load on the network (*the number of packets sent to the network*) is greater than the capacity of the network (*the number of packets a network can handle*).
- Congestion control refers to the mechanisms and techniques to control the congestion and keep the load below the capacity.
- Congestion in a network or internetwork occurs because routers and switches have queues- buffers that hold the packets before and after processing.

Network Congestion

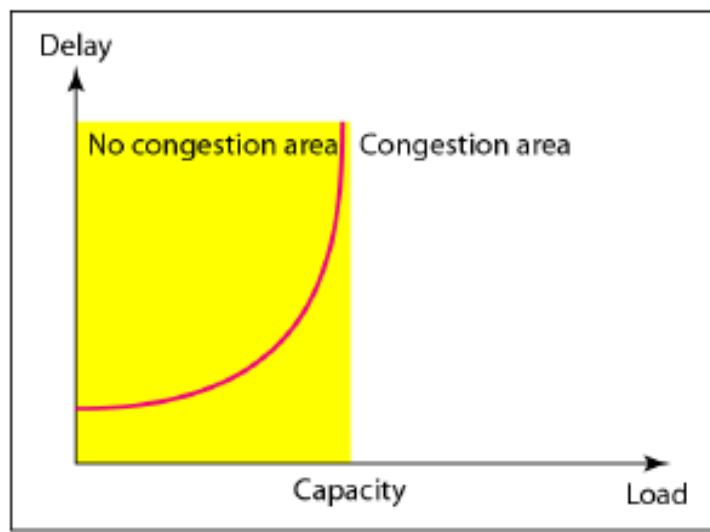
- A **router**, for example, has an **input queue** and an **output queue** for each **interface**. When a packet arrives at the **incoming interface**, it undergoes **three steps** before departing.
 1. The **packet** is put at the end of the **input queue** while waiting to be checked.
 2. The **processing module** of the router removes the packet from the input queue once it reaches the front of the queue and uses its **routing table** and the **destination address** to find the **route**.
 3. The **packet** is put in the appropriate **output queue** and **waits its turn** to be sent.
- If the **packet arrival rate** is **higher than the packet processing rate**, the **input queues** become **longer and longer**.
- If the **packet departure rate** is **less than the packet processing rate**, the **output queues** become **longer and longer**.

Queues in a Router

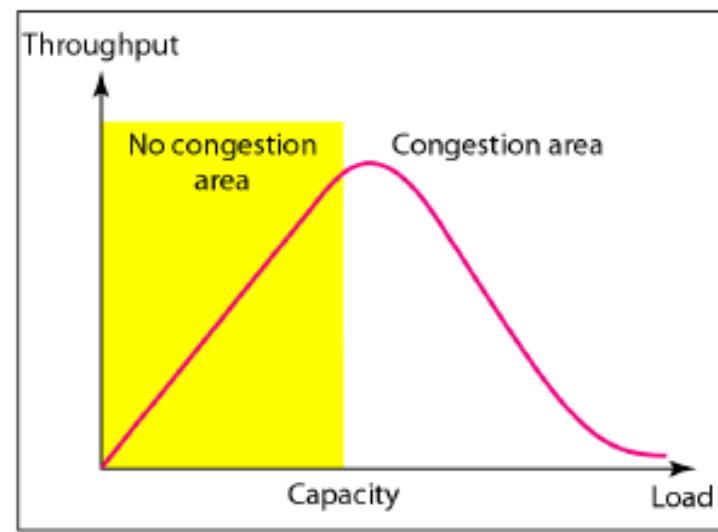


Network Performance

- **Congestion control** involves **two factors** that measure the **performance** of a network: **delay** and **throughput**.
- **Figure** below shows these two performance measures as **function** of **load**.



a. Delay as a function of load



b. Throughput as a function of load

Network Performance: Delay versus Load

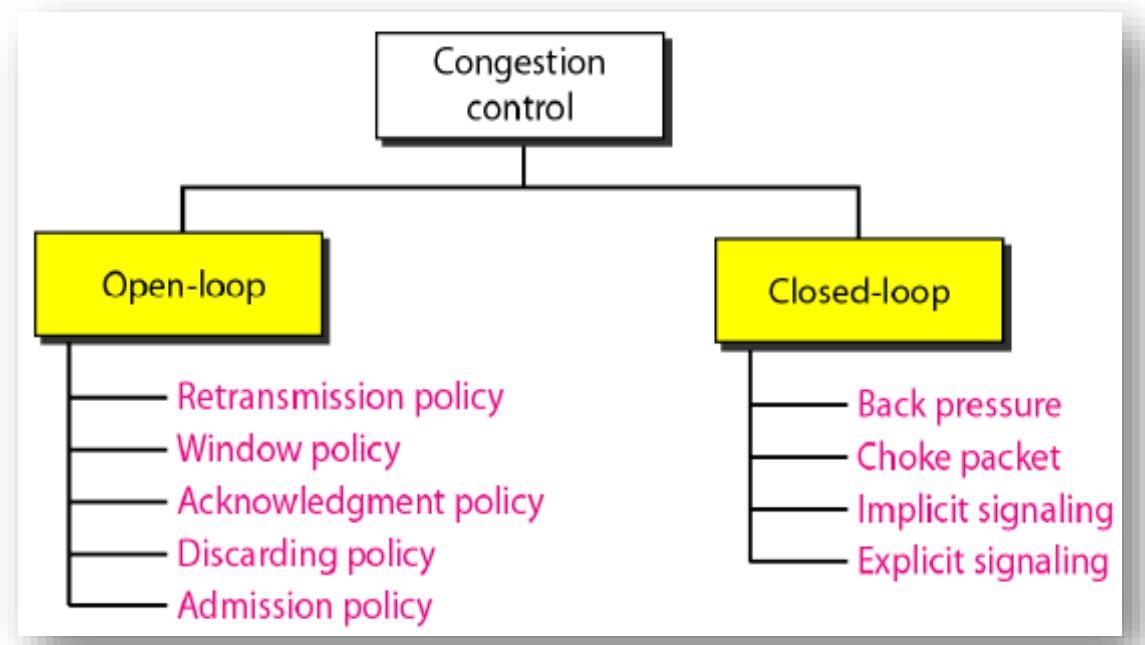
- When the **load** is much **less than** the **network capacity**, the **delay** is at a **minimum**.
- This **minimum delay** is composed of **propagation delay** and **processing delay**, both of which are **negligible**.
- However, when the **load reaches** the **network capacity**, the **delay increases sharply** because we now need to add the **waiting time in the queues** (for all routers in the path) to the **total delay**.
- The **delay becomes infinite** when the **load** is greater than the **capacity**.
- Due to **infinite delay**; the **queues** become longer and longer.
- **Delay** has a **negative effect** on the **load** and consequently the **congestion occurs**.
- When a **packet is delayed**, the **source, not receiving the acknowledgment, retransmits the packet**, which makes the **delay, and the congestion, worse**.

Network Performance: Throughput versus Load

- **Throughput** in a **network** is defined as **the number of packets** passing through the **network in a unit of time**.
- When the **load** is **below the network capacity**, the **throughput** increases **proportionally with the load**.
- We expect the **throughput** to remain constant after **the load reaches the capacity**, but instead the **throughput declines sharply**. The reason is the **discarding of packets by the routers**.
- When the **load exceeds the capacity**, the **queues become full** and the **routers have to discard some packets**.
- **Discarding packets** does **not reduce the number of packets in the network** because the **sources retransmit** the packets, using **time-out mechanisms**, when the packets do not reach the **destinations**.

Congestion Control

- Congestion control refers to techniques and mechanisms that can either prevent congestion, before it happens, or remove congestion, after it has happened.
- In general, we can divide congestion control mechanisms into two broad categories:
 - Open-loop congestion control (prevention)
 - Closed-loop congestion control (removal)



Open-Loop Congestion Control

- In **Open-loop congestion control**, policies are applied to *prevent congestion before it happens.*
- In these mechanisms, **congestion control** is **handled** by either the **source** or the **destination**.

Open-Loop Congestion Control

1. Retransmission Policy

- If the **sender** feels that a **sent packet** is **lost** or **corrupted**, the packet needs to be **retransmitted**.
- **Retransmission** in general **may increase** congestion in the network as it **increases** the **network load**.
- The **retransmission policy** and the **retransmission timers** must be designed to **optimize efficiency (minimize load)** and at the same time **prevent** congestion.
- For **example**, the **retransmission policy** used by **TCP** is designed to prevent or alleviate congestion.

Open-Loop Congestion Control

2. Window Policy

- The **type of window** at the **sender** may also **affect congestion**.
- The **Selective Repeat window** is **better than the Go-Back-N window** for **congestion control**.
- In the **Go-Back-N window**, when the timer for a packet times out, **several packets may be resent**, although some may have arrived safe and sound at the receiver.
- This **duplication** may make the **congestion worse**.
- The **Selective Repeat window**, on the other hand, tries to **send the specific packets** that have been **lost or corrupted**.

Open-Loop Congestion Control

3. Acknowledgment Policy

- The **acknowledgment policy** imposed by the receiver may also **affect congestion**.
- If the **receiver** does **not acknowledge every packet** it receives, it may **slow down** the **sender** and help **prevent congestion**.
- A **receiver** may decide to **acknowledge** only **N packets** at a time which is called **cumulative acknowledgement policy**.
- We need to know that the **acknowledgments** are also part of the **load** in a **network**.
- Sending **fewer acknowledgments** means imposing **fewer loads** on the **network**.

Open-Loop Congestion Control

4. Discarding Policy

- A good discarding policy by the routers may prevent congestion and at the same time may not harm the integrity of the transmission.
- For example, in audio transmission, if the policy is to discard less sensitive packets when congestion is likely to happen, the quality of sound is still preserved and congestion is prevented or alleviated.

5. Admission Policy

- An admission policy can also prevent congestion in virtual-circuit networks.
- Routers first check the resource requirement of a flow before admitting it to the network.
- A router can deny establishing a virtual circuit connection if there is congestion in the network or if there is a possibility of future congestion.

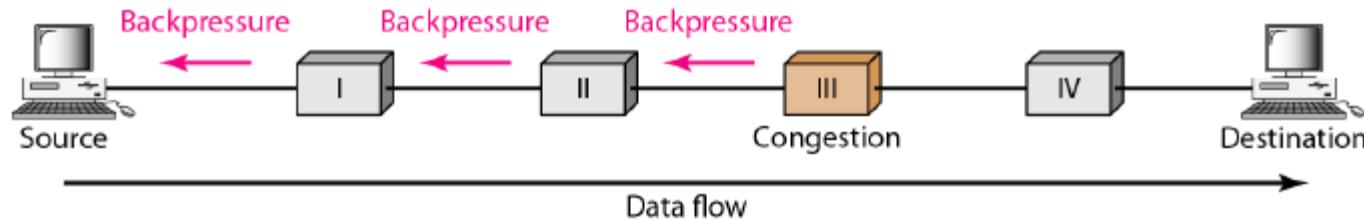
Closed-Loop Congestion Control

Closed-loop congestion control mechanisms try to alleviate congestion after it happens.

1. Backpressure

- The technique of **backpressure** refers to a **congestion control** mechanism in which a **congested node stops receiving data** from the **immediate upstream node** or nodes.
- This may cause the **upstream node or nodes to become congested**, and they, in turn, **reject data** from their **upstream nodes or nodes**. And so on.
- **Backpressure** is a **node-to-node congestion control** that starts with a node and propagates, in the opposite direction of data flow, to the source.
- The **backpressure technique** can be applied only to **virtual circuit networks**, in which **each node knows the upstream node** from which a flow of data is coming.

Closed-Loop Congestion Control



- **Node III** has more input data than it can handle.
- It drops some packets in its **input buffer** and informs **node II** to slow down.
- **Node II**, in turn, may be **congested** because it is **slowing down** the **output flow** of data.
- If **node II** is **congested**, it informs **node I** to **slow down**, which in turn may **create** congestion.
- If so, **node I** informs the **source** of data to **slow down**.
- This, in time, **alleviates** the **congestion**.
- Note that the **pressure on node III** is **moved backward** to the **source** to remove the **congestion**.

Closed-Loop Congestion Control

- Backpressure technique was implemented in the first virtual-circuit network, X.25.
- The technique cannot be implemented in a datagram network because in this type of network, a node (router) does not have the slightest knowledge of the upstream router.

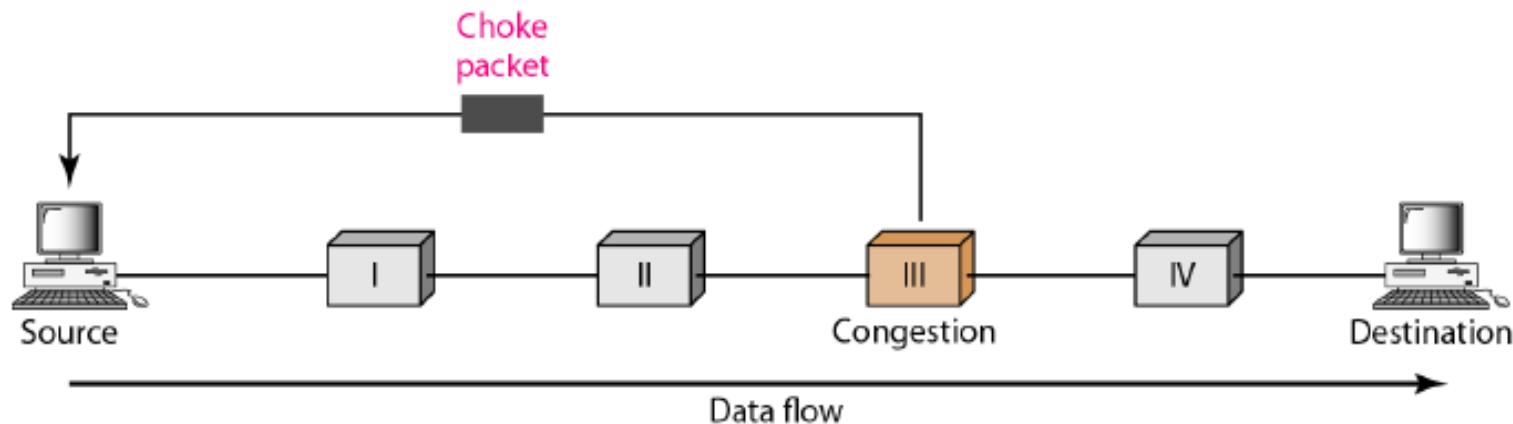
Closed-Loop Congestion Control

2. Choke Packet

- A **choke packet** is a **packet sent by a node** to the **source** to **inform** it of **congestion**.
- Note the **difference** between the **backpressure** and **choke packet** methods.
- *In **backpressure**, the **warning** is from **one node** to its **upstream node**, although the **warning** may **eventually reach** the **source station**.*
- In the **choke packet** method, the **warning** is from the **router**, which has **encountered congestion**, to the **source station directly**.
- The **intermediate nodes** through which the **packet** has **traveled** are **not warned**.
- We have seen an **example** of this **type** of control in **ICMP**.

Choke packet

- When a **router** in the **Internet** is **overwhelmed** with IP datagrams, it may **discard** some of them; but it **informs** the **source host**, using a **source quench ICMP message**.
- The **warning message** goes **directly to the source station**; the **intermediate routers** does not take any action.



Closed-Loop Congestion Control

3. Implicit Signaling

- In **implicit signaling**, there is **no communication between the congested node or nodes and the source**.
- The **source guesses** that **there** is a **congestion somewhere** in the **network** from other **symptoms**.
- For **example**, when a **source** sends **several packets** and there is **no acknowledgment** for a while, one **assumption** is that the **network** is **congested**.
- The **delay** in receiving an **acknowledgment** is **interpreted** as **congestion** in the **network**; the **source** should **slow down**.
- **TCP** uses such **signaling** for **congestion control**.

Closed-Loop Congestion Control

4. Explicit Signaling

- The node that experiences **congestion** can **explicitly** send a **signal** to the **source or destination**.
- The **explicit signaling** method, however, is **different** from the **choke packet method**.
- In the **choke packet** method, a **separate packet** is used for this purpose; in the **explicit signaling method**, the **signal** is **included** in the **packets that carry data**.
- **Explicit signaling**, as used in **Frame Relay congestion control**, can occur in either the **forward** or the **backward** direction.

Closed-Loop Congestion Control

4.1. Backward Signaling

- A **bit** can be **set** in a **packet moving** in the **direction opposite** to the **congestion**.
- This **bit** can **warn the source** that there is **congestion** and that it **needs to slow down** to avoid the discarding of packets.

4.2. Forward Signaling

- A **bit** can be **set** in a **packet moving** in the **direction of the congestion**.
- This **bit** can **warn the destination** that there is **congestion**.
- The **receiver** in this **case** can use **policies**, such as **slowing down** the **acknowledgments**, to **alleviate the congestion**.

Congestion Control in TCP

- TCP uses **congestion control** to **avoid congestion** or **alleviate congestion** in the network.

Congestion Window

- The **sender window size** is determined by the available buffer space in the receiver (*rwnd*).
- In addition to the receiver, the **network** is a second entity that determines the size of the sender's window.
- Thus, *the sender's window size is determined not only by the receiver but also by congestion in the network.*
- The **sender has two pieces of information**: the **receiver-advertised window size** and the **congestion window size**.
- The actual size of the window is the **minimum of these two**.
- i.e. **Sender's window size= minimum (rwnd, cwnd)**

TCP Congestion Policy

- TCP's general policy for handling congestion is based on **three phases**:
 1. *Slow start(SS)*
 2. *Congestion avoidance(CA)*
 3. *Congestion detection(CD)*
- In the **slow-start phase**, the sender starts with a **very slow rate of transmission**, but **increases the rate rapidly** to reach a **threshold**.
- When the **threshold** is reached, the data rate is reduced to avoid congestion.
- Finally if **congestion is detected**, the sender goes back to the **slow-start or congestion avoidance phase** based on **how the congestion is detected**.

1. Slow Start: Exponential Increase

- One of the **algorithms** used in TCP congestion control is called ***slow start***.
- This **algorithm** is based on the idea that the **size** of the **congestion window (*cwnd*)** starts with one maximum segment size (**MSS**).
- The **MSS** is **determined** during **connection establishment** phase.
- The **size** of the **window** increases **one MSS** each time an **acknowledgment** is received.
- As the name implies, the **window starts slowly, but grows exponentially**.
- We have assumed that ***rwnd*** is much higher than ***cwnd***, so that the **sender window size always equals *cwnd***.
- We have assumed that each **segment** is **acknowledged** individually.

1. Slow Start: Exponential Increase

- The **Sender starts** with **cwnd =1 MSS**.
- This means that the **sender** can send only **one segment**.
- After **receipt** of the **acknowledgment** for **segment 1**, the size of the **congestion window** is **increased by 1**, which means that **cwnd** is now **2**.
- Now two more **segments** can be **sent**.
- When each **acknowledgment** is received, the size of the **window** is increased by **1 MSS**.
- When all **seven segments** are **acknowledged**, **cwnd = 8**.

Slow Start: Exponential Increase

- If we look at the **size of *cwnd*** in terms of rounds (acknowledgment of the whole window of segments), we find that the rate is exponential as shown below:

Start *cwnd*=1

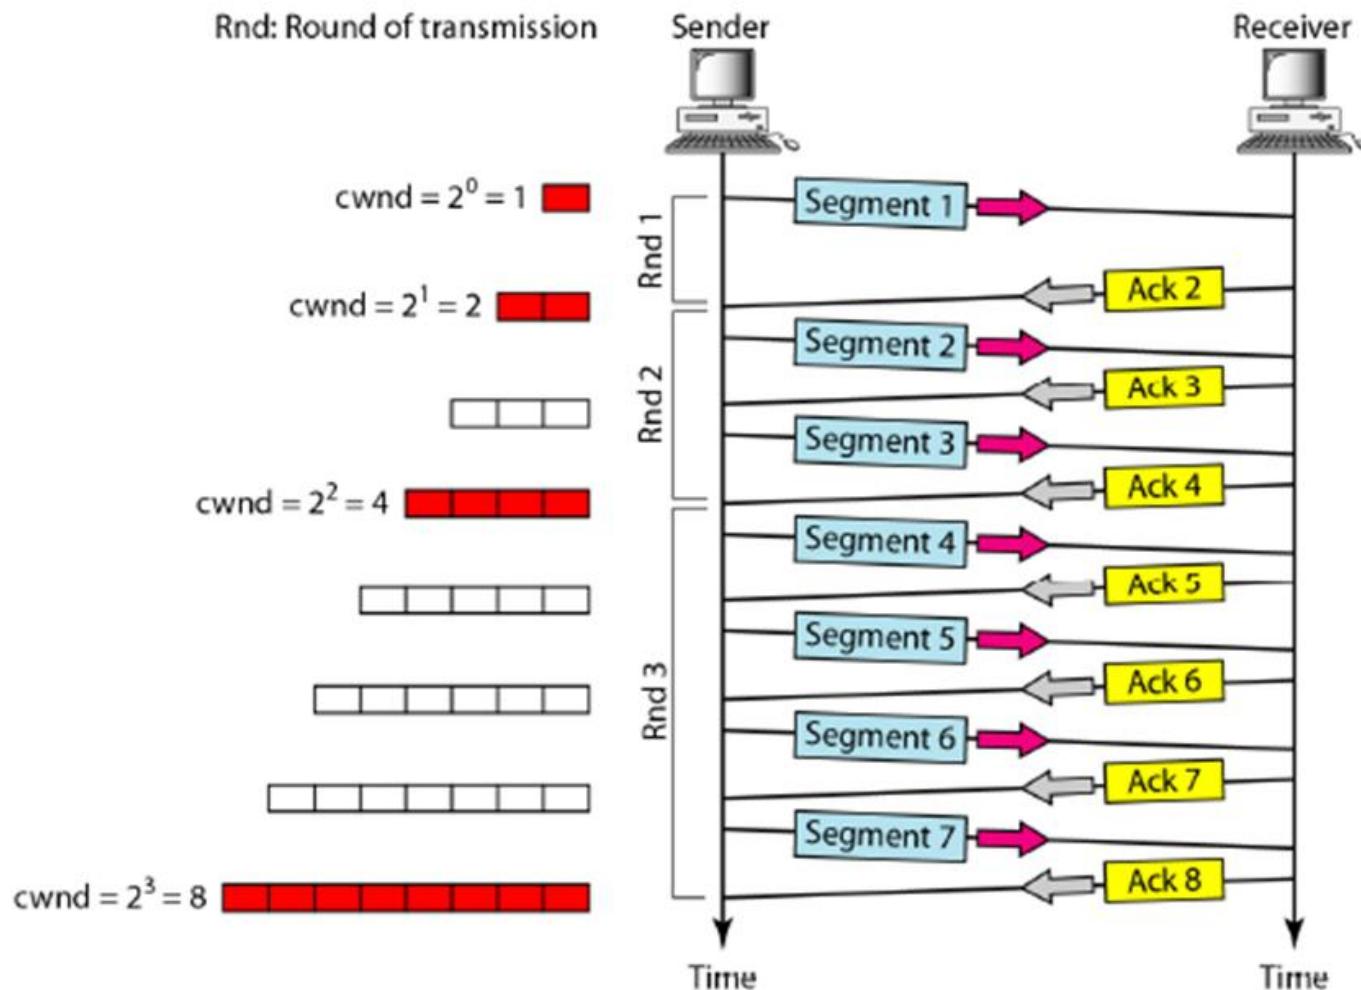
After round 1 *cwnd*= $2^1 = 2$

After round 2 *cwnd*= $2^2 = 4$

After round 3 *cwnd*= $2^3 = 8$

- We need to mention that **if there is delayed ACKs, the increase in the size of the window is less than power of 2.**
- **Slow start cannot continue indefinitely.** There must be a **threshold to stop this phase.**
- The **sender keeps track of a variable named *ssthresh*** (slow-start threshold).
- When the **size of window in bytes reaches this threshold, slow start stops and the next phase starts.**
- In most implementations the value of ***ssthresh*** is **65,535 bytes.**

Slow Start: Exponential Increase



Congestion Avoidance: Additive Increase

- **Congestion avoidance** undergoes an **additive increase** instead of an exponential one.
- When the size of the **congestion window reaches the slow-start threshold**, the **slow-start phase stops and the additive phase begins**.
- In this **algorithm**, each time the whole window of segments is acknowledged (one round), the **size of the congestion window is increased by 1**.
- **Congestion avoidance** algorithm usually starts when the size of the window is much greater than 1.

Congestion Avoidance: Additive Increase

- In this case, after the sender has **received acknowledgments** for a **complete window size of segments**, the size of the window is **increased by one segment**. If we look at the size of *cwnd* in terms of rounds, we find that the rate is additive as shown below:

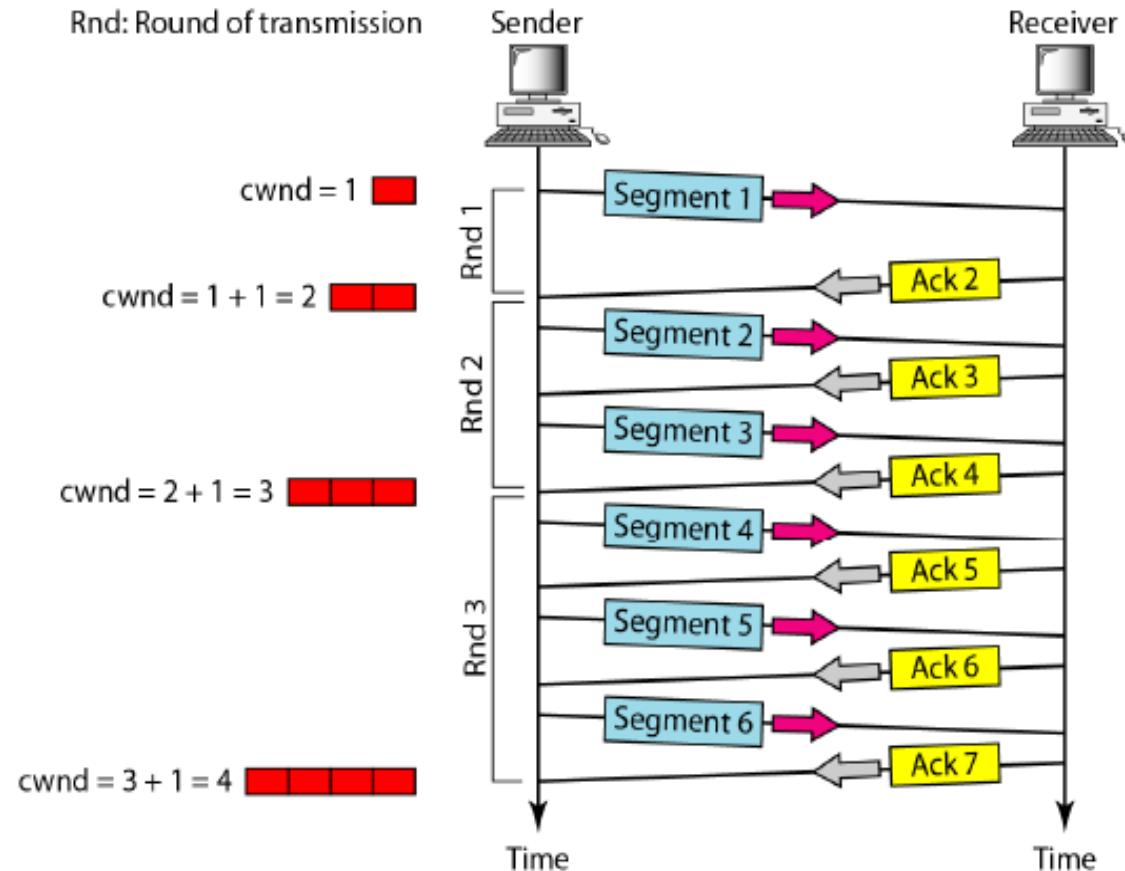
Start 1 *cwnd*=1

After round 1 *cwnd*= 1+ 1 =2

After round 2 *cwnd*=2+ 1 =3

After round 3 *cwnd*=3+ 1 =4

Congestion avoidance, additive increase



Congestion Detection: Multiplicative Decrease

- If congestion occurs, the **congestion window size must be decreased**.
- The only way the sender can guess that **congestion has occurred** is by the need to **retransmit a segment**.
- However, **retransmission can occur in one of two cases**: when a **timer times out** or when **three duplicate ACKs are received**.
- In both cases, the size of the **threshold is dropped to one-half**, a **multiplicative decrease**.
- Most **TCP implementations** have **two cases**:

Congestion Detection: Multiplicative Decrease

1. If a **time-out** occurs, there is a **stronger possibility** of congestion; a segment has probably been **dropped in the network**, and there is **no news about the sent segments**.
 - In this case **TCP reacts strongly**:
 - a. It sets the value of the **threshold** to **one-half** of the **current window size**.
 - b. It sets **cwnd** to the **size of one segment**.
 - c. It **starts** the **slow-start phase** again.

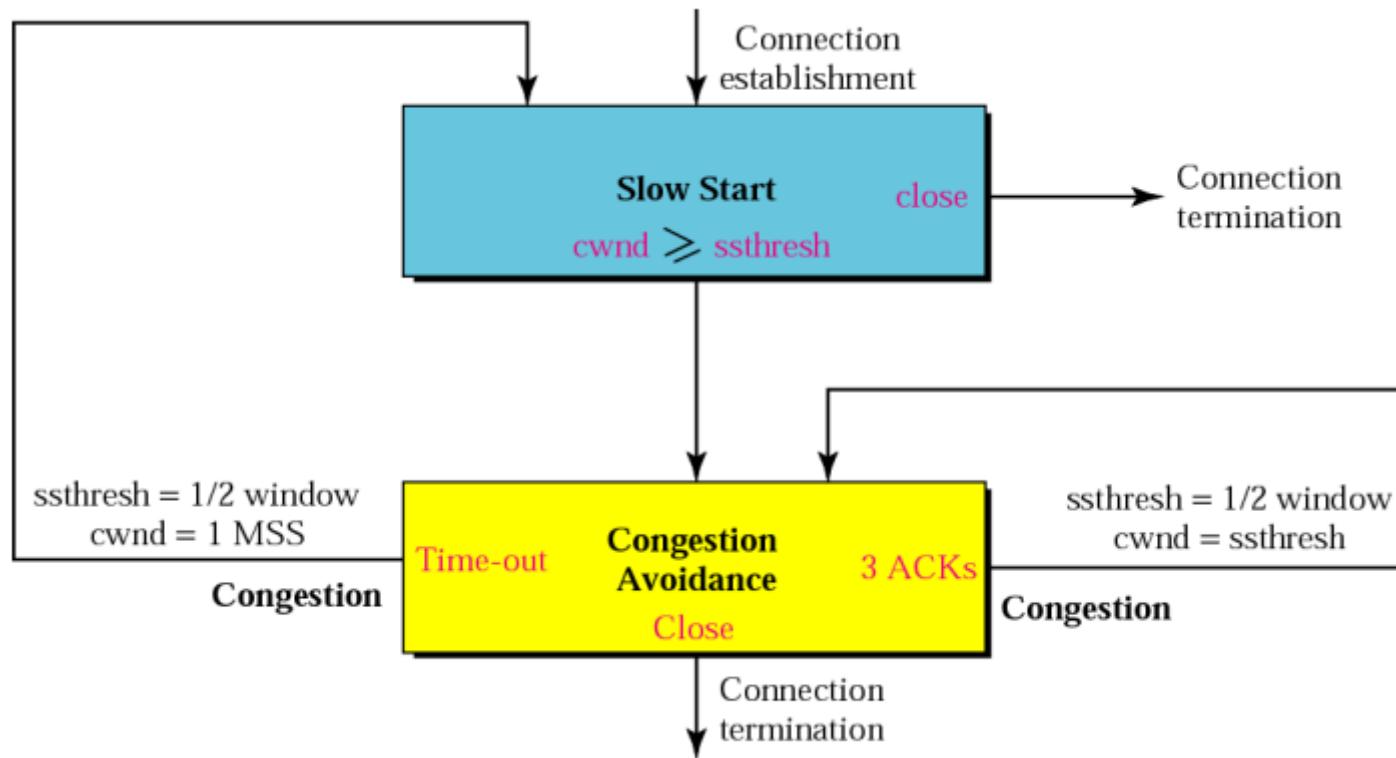
Congestion Detection: Multiplicative Decrease

2. If **three DUP ACKs** are received, there is a **weaker possibility** of congestion; a segment **may have been dropped**, but **some segments after that may have arrived safely since three ACKs are received**. This is called **fast transmission** and **fast recovery**.

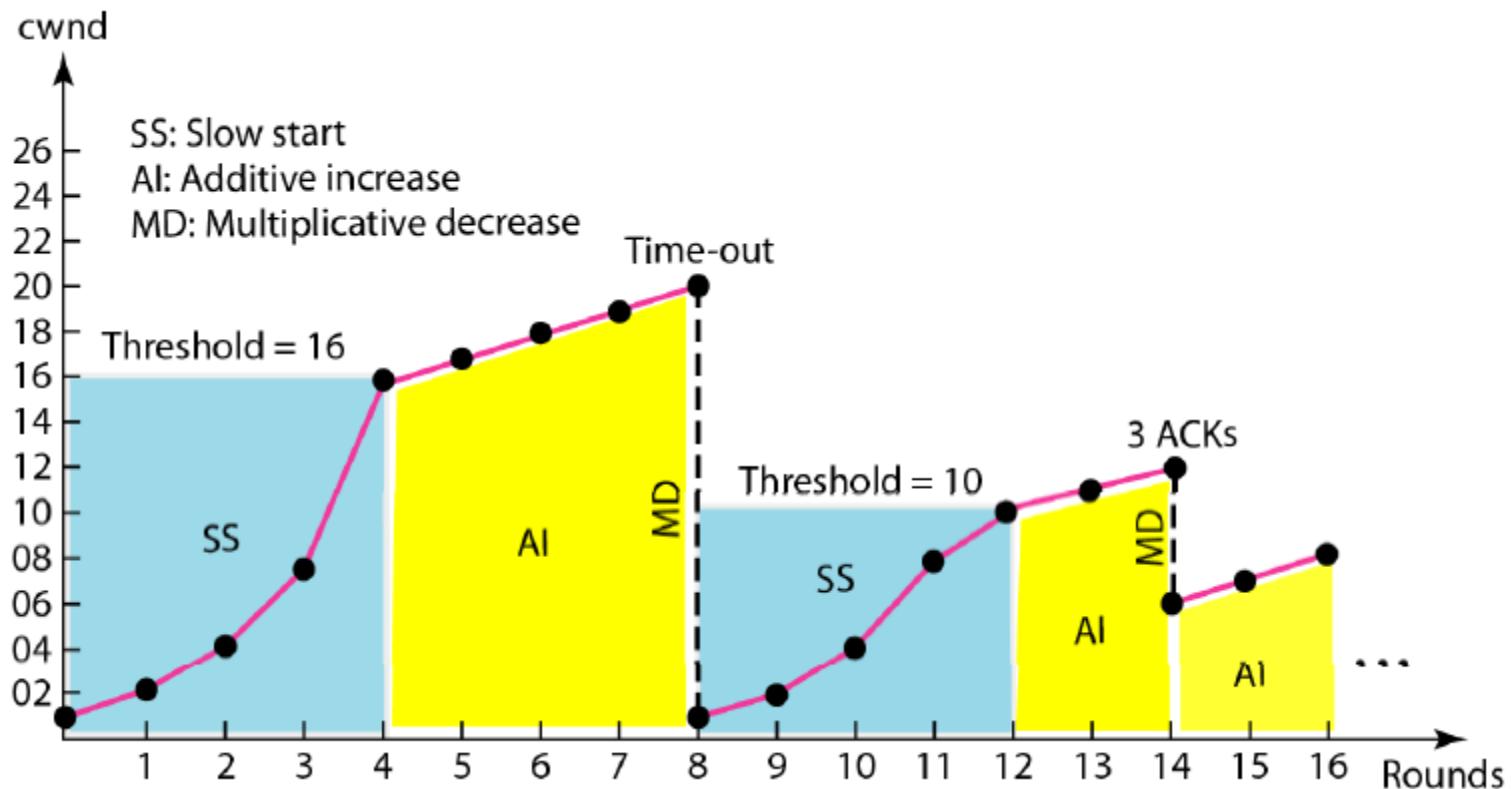
In this case, TCP has a **weaker reaction**:

- a. It sets the value of the **threshold** to **one-half** of the **current window size**.
- b. It sets **cwnd** to the **value** of the **threshold**
- c. It starts the **congestion avoidance phase**.

TCP Congestion Policy Summary



TCP Congestion Control Example



- We assume that the **maximum window size** is **32 segments**. The **threshold** is set to **16 segments** (one-half of the maximum window size).
- In the ***slow-start* phase** the window size **starts from 1** and grows **exponentially** until it reaches the threshold.

TCP Congestion Control Example

- After it reaches the **threshold**, the **congestion avoidance (additive increase)** procedure allows the window size to increase linearly until a timeout occurs or the maximum window size is reached. In Figure, the **time-out occurs** when the **window size is 20**.
- At this moment, the **multiplicative decrease** procedure takes over and reduces the threshold to one-half of the previous window size. The previous window size was 20 when the time-out happened so the new threshold is now 10.
- TCP moves to **slow start again** and starts with a window size of 1, and TCP moves to additive increase when the new threshold is reached.
- When the window size is 12, a **three-DUPACKs** event happens. The **multiplicative decrease** procedure takes over again.
- The **threshold** is set to **6** and **TCP** goes to the **additive increase phase** this time.
- It remains in this phase until another time-out or another three ACKs happen.