

# **Lecture 7.1**

## **Application Layer: Domain Name System(DNS)**


**Dr. Vandana Kushwaha**

Department of Computer Science  
Institute of Science, BHU, Varanasi

# Introduction

- There are several **applications** in the **application layer** of the **Internet model** that follow the **client/server paradigm**.
- The **client/server programs** can be divided into **two categories**: those that can be directly used by the **user**, such as **e-mail**, and those that **support other application programs**.
- The **Domain Name System (DNS)** is a **supporting program** that is used by other programs such as **e-mail**.
- A user of an **e-mail** program may know the **e-mail address** of the **recipient**; however, the **IP protocol** needs the **IP address**.
- The **DNS client** program sends a **request** to a **DNS server** to **map** the **e-mail address** to the **corresponding IP address**.
- **Figure** on next slide shows an **example** of how a **DNS client/server** program can support an **e-mail program** to find the **IP address** of an **e-mail recipient**.

# Example of using the DNS service



# Domain Name System


- To identify an entity, TCP/IP protocols use the IP address, which uniquely identifies the connection of a host to the Internet.
- However, people prefer to use names instead of numeric addresses.
- Therefore, we need a system that can map a name to an address or an address to a name.
- When the Internet was small, mapping was done by using a host file.
- The host file had only two columns: name and address.
- Every host could store the host file on its disk and update it periodically from a master host file.
- When a program or a user wanted to map a name to an address, the host consulted the host file and found the mapping.
- Today, however, it is impossible to have one single host file to relate every address with a name and vice versa.

# Domain Name System

- The **host file** would be **too large** to store in **every host**.
- In addition, it would be **impossible** to **update all** the **host files** **every time** there was a **change**.
- One **solution** would be to **store** the **entire host file** in a **single computer** and allow access to this centralized information to every computer that needs **mapping**.
- But it would **create a huge amount of traffic** on the Internet.
- Another **solution**, the one used today, is to **divide this huge amount of information** into **smaller parts** and **store each part** on a different computer.
- In this **method**, the **host** that needs **mapping** can contact the **closest computer holding** the needed information.
- This **method** is used by the **Domain Name System (DNS)**.

# DOMAIN NAME SPACE

- A **domain name space** was designed using **hierarchical name space**.
- In this **design** the **names are defined** in an **inverted-tree structure** with the **root** at the **top**.
- The **tree** can have only **128 levels**: **level 0 (root)** to **level 127**.



# DOMAIN NAME SPACE


## Label

- Each **node** in the **tree** has a **label**, which is a **string** with a maximum of **63 characters**.
- The **root label** is a **null string** (empty string).
- **DNS** requires that **children of a node** (nodes that branch from the same node) have **different labels**, which guarantees the **uniqueness** of the **domain names**.

## Domain Name

- Each **node** in the **tree** has a **domain name**.
- A **full domain name** is a **sequence of labels separated by dots (.)**
- The **domain names** are always **read** from the **node up to the root**.
- The **last label** is the **label of the root (null)**.
- This means that a **full domain name** always **ends in a null label**, which means the **last character** is a **dot** because the **null string** is nothing.

# Domain Name




# Distribution of Name Space

- The **information** contained in the **domain name space** must be **stored somewhere**.
- However, it is very **inefficient** and also **unreliable** to have just **one computer** store such a **huge amount of information**.
- It is **inefficient** because responding to requests from **all over the world** places a **heavy load** on the system.
- It is **unreliable** because any **failure** makes the **data inaccessible**.

# Hierarchy of Name Servers


- The **solution** to these **problems** is to **distribute** the **information** among **many computers** called **DNS servers**.
- One way to do this is to **divide the whole space** into **many domains** based on the **first level**.
- In other words, we let the **root** stand alone and create as **many domains (subtrees)** as there are **first-level nodes**.
- Because a **domain** created in this way could be **very large**, DNS allows domains to be **divided further** into **smaller domains (subdomains)**.
- Each **server** can be **responsible** (authoritative) for a **small domain**.
- In other words, we have a **hierarchy of servers** in the same way that we have a **hierarchy of names**.

# Hierarchy of name servers




# DNS IN THE INTERNET

- DNS is a protocol that can be used in different platforms.
- In the Internet, the domain name space (tree) is divided into three different sections: *Generic domains*, *Country domains*, and the *Inverse domain*.



# Generic Domains

- The **Generic domains** define **registered hosts** according to their **generic behavior**.
- **Each node** in the **tree** defines a **domain**, which is an **index** to the **domain name space database**.



# Generic Domains


- First level in the **generic domains** section allows **14 possible labels**.
- These **labels** describe the **organization types** as listed in **Table** below.

| <i>Label</i> | <i>Description</i>                        |
|--------------|---|
| aero         | Airlines and aerospace companies          |
| biz          | Businesses or firms (similar to "com")    |
| com          | Commercial organizations                  |
| coop         | Cooperative business organizations        |
| edu          | Educational institutions                  |
| gov          | Government institutions                   |
| info         | Information service providers             |
| int          | International organizations               |
| mil          | Military groups                           |
| museum       | Museums and other nonprofit organizations |
| name         | Personal names (individuals)              |
| net          | Network support centers                   |
| org          | Nonprofit organizations                   |

# Country Domains

- The **country domains** section uses **two-character country abbreviations** (e.g., **us** for United States).
- **Second labels** can be **organizational**, or they can be more specific, national designations.
- The **United States**, for **example**, uses **state abbreviations** as a **subdivision** of us (e.g., **ca.us.**).
- The address ***anza.cup.ca.us*** can be translated to **De Anza College in Cupertino, California**, in the **United States**.

# Country Domains



# Inverse Domain

- The **inverse domain** is used to **map** an **address** to a **name**.
- This may happen, for **example**, when a **server** has received a **request** from a **client** to do a **task**.
- Although the **server** has a **file** that contains a **name list of authorized clients**.
- The **server** asks its **resolver** to **send** a **query** to the **DNS server** to **map an address to a name** to determine if the **client is on the authorized list**.


# Address Resolution

- **Mapping** a **name** to an **address** or an **address** to a **name** is called ***name-address resolution***.

## Resolver

- DNS is designed as a **client/server application**.
- A **host** that needs to **map** an **address** to a **name** or a **name** to an **address** calls a **DNS client** called a **resolver**.
- The **resolver** accesses the **closest DNS server** with a **mapping request**.
- If the **server** has the **information**, it **satisfies** the **resolver**; otherwise, it either **refers** the **resolver** to **other servers** or asks **other servers** to provide the **information**.
- After the **resolver receives** the **mapping**, it **interprets** the **response** to see if it is a **real resolution** or an **error**, and **finally delivers** the **result** to the **process** that **requested it**.

# DNS Resolver




# Address Resolution

## Recursive Resolution

- The **client (resolver)** can ask for a **recursive answer** from a **name server**.
- This means that the **resolver expects the server to supply the final answer**.
- If the **server** is the authority for the **domain name**, it **checks** its **database** and **responds**.
- If the **server** is **not** the authority, it **sends** the **request** to another server (the parent usually) and **waits** for the **response**.
- If the **parent** is the **authority**, it **responds**; otherwise, it **sends** the **query** to yet **another server**.
- When the **query** is **finally resolved**, the **response travels back** until it **finally reaches** the **requesting client**.


# Recursive Resolution



# Iterative Resolution

- If the **client** does not ask for a **recursive answer**, the **mapping** can be done **iteratively**.
- If the **server** is an **authority** for the **name**, it **sends** the **answer**.
- If it is **not**, it **returns** (to the client) the **IP address** of the **server** that it thinks can **resolve the query**.
- The **client** is responsible for **repeating** the **query** to this **second server**.
- If the **newly addressed server** can **resolve** the **problem**, it **answers** the **query** with the **IP address**; otherwise, it **returns** the **IP address** of a **new server** to the **client**.
- Now the **client** must **repeat the query** to the **third server**.
- This **process** is called **iterative resolution** because the client **repeats the same query** to **multiple servers**.

# Iterative Resolution



# Caching

- Each time a **server** receives a query for a name that is **not** in its domain, it needs to search its **database** for a server IP address.
- **Reduction** of this **search time** would **increase efficiency**.
- **DNS** handles this with a **mechanism** called **caching**.
- When a **server** asks for a **mapping** from another server and **receives the response**, it **stores this information** in its **cache memory** before **sending** it to the **client**.
- If the **same** or **another client** asks for the **same mapping**, it can **check** its **cache memory** and **solve** the **problem**.
- **Caching** speeds up resolution, but it can also be problematic.
- If a **server caches** a **mapping** for a **long time**, it may send an **outdated mapping** to the **client**.
- To counter this, **two techniques** are used.

# Caching

- First, the authoritative server always adds information to the mapping called *time-to-live (TTL)*.
- It defines the time in seconds that the receiving server can cache the information.
- After that time, the mapping is invalid and any query must be sent again to the authoritative server.
- Second, DNS requires that each server keep a TTL counter for each mapping it caches.
- The cache memory must be searched periodically, and those mappings with an expired TTL must be purged.

# ENCAPSULATION

- DNS can use either **UDP or TCP**. In both cases the **well-known port** used by the server is **port 53**.
- UDP is used when the **size** of the **response message** is **less than 512 bytes** because most **UDP packages** have a **512-byte packet size limit**.
- If the **size** of the **response message** is **more than 512 bytes**, a **TCP connection** is used.
- In that case, one of two scenarios can occur:

## Case 1:

- If the **resolver** has **prior knowledge** that the **size** of the **response message** is **more than 512 bytes**, it uses the **TCP connection**.

# ENCAPSULATION

## Case 2:

- If the **resolver** does not know the **size** of the **response message**, it can use the **UDP port**.
- However, if the **size** of the **response message** is **more than 512 bytes**, the **server truncates the message** and **turns on the TC bit**.
- The **resolver** now **opens a TCP connection** and **repeats the request** to get a **full response** from the **server**.
- Note: **TC Truncated : 1 bit in DNS header**, response **too large for UDP (1)**.