

Lecture 13

Query Processing and Optimization

Dr. Vandana Kushwaha

Department of Computer Science
Institute of Science, BHU, Varanasi

Query Processing

- A **database query** expressed in a **high-level query language** such as **SQL** must first be *scanned, parsed*, and *validated*.
- The **scanner** identifies the **query tokens**—such as **SQL keywords**, **attribute names**, and **relation names** that appear in the text of the **query**.
- The **parser** checks the **query syntax** to **determine** whether it is **formulated according** to the **syntax rules** (rules of grammar) of the **query language**.
- The **query** must also be **validated** by checking that **all attribute** and **relation names** are **valid** and **semantically meaningful names** in the **schema** of the particular **database** being queried.
- An **internal representation** of the **query** is then **created**, usually as a **tree data structure** called a **query tree**.

Query Optimization

- The **internal representation** of a **query** i.e. **query tree** is now passed to the **Query Optimization** module for finding the **optimal execution plan**.
- A **query** typically has **many possible execution strategies**.
- The **process** of **choosing** a **suitable execution strategies** for **processing a query** is known as **query optimization**.
- There are **two main techniques** that are **employed** during **query optimization**:
 - *Heuristic based Query Optimization*
 - *Cost based Query Optimization*


Query Optimization

- **Heuristic based Query Optimization**
 - This **technique** is based on **heuristic rules** for **ordering the operations** in a **query execution strategy**.
 - A **heuristic** is a **rule** that **works well** in **most cases** but is **not guaranteed** to **work well** in **every case**.
 - The **rules** typically **reorder the operations** in a **query tree**.

Query Optimization

- **Cost based Query Optimization**
 - This **technique** involves **systematically estimating** the **cost** of **different execution strategies** :
 - *Access cost,*
 - *Storage cost,*
 - *Computation cost,*
 - *Communication cost.*
 - And then **choosing** the **execution plan** with the **lowest cost estimate**.

Query Processing & Optimization



Query Optimization

- **Why optimizing a query?**
 - So that it is **processed efficiently**.
- **What is meant by efficiently?**
 - **Minimizing** the **cost** of **query evaluation**.
- **Who will minimize?**
 - The **system**, not the **user**.
- **Query optimization** *is facilitating a system to **construct** a **query-evaluation plan** for processing a query **efficiently**, without expecting users to **write** efficient queries.*

Query Tree


- **Tree** that represents a **relational algebra expression** corresponding to a **SQL statement** is referred as **Query Tree**.
 - **Leave nodes** represent the **base tables**.
 - **Internal nodes** represent the **relational algebra operators** applied to the **child nodes**.
 - The **query tree** is **executed** from **leaves to root**.
 - An **internal node** can be **executed** when its **operands** are **available** (children relation has been computed)
 - **Internal Node** is **replaced** by the **result of the operation** it represents.
 - **Root node** is **executed last**, and is **replaced** by the **result** of the **entire tree**.

Canonical query tree

- Construct the **Canonical(initial) query tree** as follows:
 - **Cartesian product** of the **FROM-tables**
 - **Select** with **WHERE-condition**
 - **Project** to the **SELECT-attributes**

- **Example:**

SELECT attributes
FROM A, B, C
WHERE condition




Canonical Query Tree

- *Example: List the last name of the employees born after 1957 who work on a project named "Aquarius".*

SELECT *E*.LNAME **FROM** EMPLOYEE *E*, WORKS_ON *W*, PROJECT *P*

WHERE *P*.PNAME = 'Aquarius' **AND** *P*.PNUMBER = *W*.PNO **AND** *W*.ESSN = *E*.SSN
AND *E*.BDATE > '1957-12-31'



Heuristics based Query Optimization


- **Heuristic rules** are used to **modify** the **internal representation** of a **query**—which is usually in the form of a **query tree** data structure—to **improve** its **expected performance**.
- The **scanner** and **parser** of an **SQL query** first generate a **data structure** that corresponds to an ***initial query representation***, which is then **optimized** according to **heuristic rules**.
- This leads to an ***optimized query representation***, which corresponds to the **query execution strategy**.
- Following that, a **query execution plan** is **generated** to **execute** groups of **operations** based on the **access paths** available on the **files involved** in the **query**.

Heuristics based Query Optimization

- One of the **main heuristic rules** is :
- *Apply **SELECT** and **PROJECT** operations before applying the **JOIN** or other **binary operations**.*
- Because the **size** of the **file** resulting from a **binary operation**—such as **JOIN**—is usually a **multiplicative function** of the **sizes** of the **input files**.
- The **SELECT** and **PROJECT** operations **reduce the size** of a **file** and hence should be **applied before a join** or other **binary operation**.
- Thus in a **Query tree** we should **try to move** **SELECT** and **PROJECT** operations **down**.


Heuristics based Query Optimization

- **Example:** Move *SELECT* down:




Heuristics based Query Optimization

- Example:** Move SELECT **down**:




Heuristics based Query Optimization

- **Example:** Apply **most restrictive Select first.**




Heuristics based Query Optimization

- **Example:** Replace **Cartesian product** with **Join**.



Heuristics based Query Optimization



- **Example:** Restrict selection results with projection.



Heuristics based Query Optimization

- **Example:** Move PROJECT down:

$\sigma_{balance < 2500}(\pi_{balance}(account))$



Rules for Transformation of Relational Expressions

- A **query** can be **expressed** in several **different ways**, with different **costs of evaluation**.
- Two **relational-algebra expressions** are said to be **equivalent** if, on every legal database instance, the **two expressions** generate the **same set of tuples**.
- **Equivalence Rules**
- An **equivalence rule** says that **expressions of two forms** are **equivalent**.
- Let $\theta_1 \theta_2 \theta_3 \dots$ Are **conditions** (predicates).
- **Rule 1:** Conjunctive selection operations can be **deconstructed** into a **sequence of individual selections**. This **transformation** is referred to as a **cascade of σ** .

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

Rules for Transformation of Relational Expressions

- **Rule 2:** Selection operations are **commutative**.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

- In practice, it is **better** and more **optimal** to **apply** that **selection first** which yields a **fewer** number of tuples (**restrictive selection**).
- This **saves time** on our **outer selection**.

Rules for Transformation of Relational Expressions

- **Rule 3:** Selections on **Cartesian Products** can be re-written as **Theta Joins**(Natural Join).

$$\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$$

- The **cross product operation** is known to be **very expensive**.
- This is because it **matches** each tuple of **E1** (total **m tuples**) with each tuple of **E2** (total **n tuples**).
- This **yields** **m*n** entries.

Rules for Transformation of Relational Expressions

- **Rule 4:** Theta Joins are **commutative**.

$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

- **Rule 5:** Join operations are **associative**.

$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

- **Joins** are all **commutative** as well as **associative**, so one must **join those two tables first** which **yield less number of entries**, and then **apply the other join**.

Rules for Transformation of Relational Expressions

- Even though **join order** doesn't affect **query results** (because joins are commutative and associative), it can **drastically affect performance** because:
 - **Smaller intermediate results** → **faster joins**
 - **Joining selective (filtered) tables first** → **fewer rows to process next**
- **Example:** Let us consider **three tables**:

Customers (1 million rows)

| customer_id | name |

Orders (10 million rows)

| order_id | customer_id |

OrderDetails (100 million rows)

| order_detail_id | order_id | product_id |

Rules for Transformation of Relational Expressions

- Suppose we want to execute following **SQL statement**:
- *SELECT * FROM Customers c*
- *JOIN Orders o ON c.customer_id = o.customer_id*
- *JOIN OrderDetails od ON o.order_id = od.order_id WHERE c.customer_id = 123;*
- **Naive join order**
- *Orders ⋈ OrderDetails → very large intermediate result*
- *then filter to customer_id = 123*
- This **join** creates **massive intermediate data** before **filtering**—it's **slow** and **memory-intensive**.

Rules for Transformation of Relational Expressions

- **Better join order:**
- If **optimizer** instead **joins Customers and Orders first (with WHERE filter):**
- *SELECT * FROM (SELECT * FROM Customers WHERE customer_id = 123) c*
- *JOIN Orders o ON c.customer_id = o.customer_id*
- *JOIN OrderDetails od ON o.order_id = od.order_id;*
- Now Only **1 customer row** is selected,
- Only **that customer's orders** are **joined** (maybe **100 rows** instead of **10 million**)
- Then join those **100 orders** with **OrderDetails** (instead of 100 million rows)
- Thus **intermediate results** are **tiny**, making **query execution way faster**.
- Joining **smaller tables first** → **reduces the number of rows passed into subsequent joins**.

Rules for Transformation of Relational Expressions

- **Rule 6:** Selection operation can be distributed.


$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$


- **Rule 7:** Projection distributes over the Theta Join.

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) = (\Pi_{L_1}(E_1)) \bowtie_{\theta} (\Pi_{L_2}(E_2))$$

Example



(a) Initial expression tree



(b) Tree after multiple transformations

Summary:Heuristics for Optimization

- *Do operations that **generate** smaller tables **first**.*
- Do **selection** as early as possible.
- Use **cascading, commutativity, and distributivity** to move selection as far down the **query tree** as possible.
- Use **associativity** to rearrange relations so the selection operation that will produce the **smallest table** will be executed **first**.
- Do **projection** early.
- Use **cascading, distributivity and commutativity** to move the projection as far **down** the **query tree** as possible.